AtliQ Contract Risk Analyzer

AI Opportunity Map: contract review from intake to human decision and approved memory (aligned to PRD v1.1)

SWIMLANES	1. Contract Intake	2. Context & Entity	3. Existing Contracts	4. Clause Analysis	5. Risk Scoring	6. Completeness	7. Brand & Counterparty	8. Negotiation	9. Final Decision	10. Learn & Memory
Steps	Receive client/partner contract or AtliQ outbound draft; identify document type and AtliQ role (supplier, buyer, partner, NDA party)	Validate AtliQ signing entity, signatory, governing law and jurisdiction against entity master data	Retrieve signed contracts and live obligations; detect conflicts	Segment the full document; extract every clause and commercial term, including clauses with no issue	Grade each material clause Low / Medium / High with explanation and confidence; rank by severity	Check required companion docs (NDA, BAA, DPA); link SOW to parent MSA; list missing information	Screen publicity/reference rights, non-disparagement and conduct clauses; summarise counterparty payment and dispute history	Surface how similar clauses were handled (standard vs exception); draft negotiation points and fallback language	Cited decision memo: proceed, negotiate, confirm internally, send to counsel, or do not sign yet	After signature, store approved obligations (effective and expiry dates) and the negotiated outcome
Human actions	Client or partner emails the contract (or AtliQ drafts first); reviewer uploads it. No internal requester	Validate signing entity; confirm geography and AtliQ role	Review potential conflicts and judge business relevance	Review flagged clauses; assess impact	Primary / secondary reviewer validates or re-grades severity; accept, negotiate or escalate	Answer routed questions (Finance, CTO or assigned reviewer); supply missing documents	Judge reputational exposure; weigh counterparty history against the commercial terms	Choose fallback position; negotiate with counterparty	KD or delegated approver decides: sign, redline, confirm, counsel or reject. Only KD signs	Approve, correct or reject AI findings before they become precedent
Digital tool + platform	File upload + tracker import; PDF/DOCX parser with OCR and parse-quality score; type and role classifier	Versioned deterministic rules for entity, jurisdiction and invoicing (rule-based, not LLM)	Hybrid lexical + semantic retrieval filtered by permission and effective date; approved obligation register	LLM structured extraction with exact citation spans; clause-coverage	LLM classification grounded in precedent and prior obligations + confidence scoring	Required-document rules + contract-family linker; missing-info routing to owners	Brand-safety classifier on versioned criteria; structured query on AtliQ payment and dispute records	LLM generation grounded in cited precedent and deal context; RAG over verified negotiation history	Review UI with coverage up front; delegation, primary/secondary reviewer, RBAC, audit trail	Obligation and precedent register, searchable by counterparty, clause, entity, date and live/expired status
Pain points	No universal template: every counterparty adds its own clauses, so every contract is a fresh ~3h read	Wrong AtliQ entity causes tax, invoicing and payment problems; client templates impose foreign law and courts	Signed contracts scattered across drives and inboxes; checking whether an old restriction expired is manual	Risk buried in long documents; clause-by-clause manual read; skimming under deadline pressure	Inconsistent judgment; 'mutual' NDAs treated as low-risk; missed indemnity risk has cost paid legal help	Missing NDA/BAA/DPA; orphan SOW with no parent MSA; unclear who must confirm scope or pricing	Brand exposure found only after signing; payment defaults and past disputes not visible at review time	No fixed negotiation policy; terms change on WhatsApp, email and calls; past concessions not recorded	CEO bottleneck: ~1 week in queue; external counsel slow, expensive, used only for big deals	Outcomes recorded inconsistently; no structured obligation register; contracts signed and forgotten
Guardrails + risks	Treat document text as untrusted data (prompt injection); poor OCR: abstain and ask for a cleaner copy	Entity mismatch blocks 'ready to sign'; unknown entity: abstain; rules cover entity, jurisdiction, documents only	Cite both new clause and prior obligation; state live / expired / superseded with date; never auto-clear	Quote exact clause with page/section; no overall verdict if clause coverage is below threshold (99%)	No commercial thresholds as policy; 'mutual' or 'standard' label never lowers scrutiny; Red blocks ready-to-sign	Deterministic checks; flag, do not certify compliance; review child SOW against parent MSA terms	Quote clause text; if no counterparty history is held, say so and never imply a clean record	Drafts for human review only; label standard / exception / unknown; no automatic redlines (deferred)	No auto-sign; AI never signs or rejects; delegation scoped, time-bound and logged; full audit history	Only human-approved outcomes are stored; AI output never becomes precedent; exceptions cannot update policy